

Project 1 - Canny Edge Detector

Project group members: 1. Aniket Bote (N-12824308)

2. Sindhu Harish (N19806874)

Instructions to run the code:

1. Clone the project and navigate to the project directory named 'canny-edge-detector'.
2. Run the below command

```
➔ python canny_edge_detector.py --input_folder input --output_folder output
```
3. The input_folder is from where the input images are read, and output_folder is where the output images are going to get saved.
4. In the given project structure input_folder is named as input and output_folder is named as output respectively.

Source Code:

1. The main file where the code execution starts:- canny_edge_detector.py

```
# Import the required libraries

import argparse
import glob
import os
import shutil

import cv2

from gaussian_smoothing import perform_gaussian_smoothing
from gradient_operation import perform_gradient_operation
from non_maxima_suppression import perform_non_maxima_suppression
from thresholding import perform_thresholding

parser = argparse.ArgumentParser()
parser.add_argument(
    '--input_folder',
    type=str,
    default='input',
    required=False,
    help='input folder with images'
)
```

```

parser.add_argument(
    '--output_folder',
    type=str,
    default='output',
    required=False,
    help='output folder to save processed
images'
)

args = parser.parse_args()

if os.path.exists(args.output_folder):
    shutil.rmtree(args.output_folder)

os.makedirs(args.output_folder)

print("Reading images from input folder")
images = glob.glob(os.path.join(args.input_folder, '*.bmp'))
for image_name in images:
    output_image_name = image_name.split('\\')[1].split('.')[0]
    img = cv2.imread(image_name, cv2.IMREAD_GRAYSCALE)

    print("Performing gaussian smoothing for image: " +
          output_image_name)
    gaussian_smooth_image =
    perform_gaussian_smoothing(args, output_image_name, img)

    print("Performing gradient smoothing for image: " +
          output_image_name)
    M, THETA = perform_gradient_operation(args, output_image_name,
    gaussian_smooth_image)

    print("Performing non-maxima suppression for image: " +
          output_image_name)
    NMS = perform_non_maxima_suppression(args, output_image_name,
    M, THETA)

    print("Performing thresholding for image: " +
          output_image_name, '\n')
    T1, T2, T3 = perform_thresholding(args, output_image_name,
    NMS)

```

2. Source Code for Gaussian Smoothing :- gaussian_smoothing.py

```
import os
import cv2
import numpy as np
from utils import Operator, apply_discrete_convolution

def perform_gaussian_smoothing(args, image_name, image):
    """
    Args:
        image : An image to on which smoothing will appear
    Returns:
        smoothed image : Smoothened image
    """
    # Apply discreet convolution with gaussian mask
    image = apply_discrete_convolution(image, Operator.gaussian_mask)

    # Normalize the image
    image = image / np.sum(Operator.gaussian_mask)

    #write image into the output folder after normalization
    cv2.imwrite(os.path.join(args.output_folder, image_name +
'_gaussian_smooth_normalized.bmp'), image)

    # Return the smoothed image
    return image
```

3. Source Code for Gradient Operation:- gradient_operation.py

```
import os
import cv2
import numpy as np
from utils import Operator, apply_discrete_convolution

def perform_gradient_operation(args, image_name, image):
    """
    Args:
        image : An image on which gradient operation will happen
    Returns:
        Magnitude : Magnitude of the gradient
        Theta      : Gradient Angle
    """
```

```

# Compute horizontal gradients
dfdx = apply_discrete_convolution(image, Operator.gx)

#Copy Image to Output folder after horizontal gradient
cv2.imwrite(os.path.join(args.output_folder, image_name +
'_Gx_normalized.bmp'), dfdx)

# Compute vertical gradients
dfdy = apply_discrete_convolution(image, Operator.gy)

#Copy Image to Output folder after vertical gradient
cv2.imwrite(os.path.join(args.output_folder, image_name +
'_Gy_normalized.bmp'), dfdy)

# Compute magnitude of the gradient
m = np.sqrt(np.square(dfdx) + np.square(dfdy))

# Normalize gradient magnitude
m = np.absolute(m) / 3

#Copy Image to Output folder with gradient magnitude value
cv2.imwrite(os.path.join(args.output_folder, image_name +
'_gradient_magnitude_normalized.bmp'), m)

# Compute gradient angle
theta = np.degrees(np.arctan2(dfdy, dfdx))

return m, theta

```

4. Source Code for non-maxima-supression:- non_maxima_supression.py

```

import os
import numpy as np
import cv2
from utils import get_positive_angle, Sector

def perform_non_maxima_suppression(args, image_name, magnitude,
gradient_angle):
    """
    Args:
        magnitude : Magnitude of the gradient
        gradient_angle : Gradient angle
    Returns:
        Magnitude : Magnitude array after non-maxima supression
    """

```

```

...
# Compute positive angles
positive_gradient_angle = get_positive_angle(gradient_angle)

# Get magnitude array shape
m_arr, n_arr = magnitude.shape

# reference pixel location during start of the process
rpi_m, rpi_n = 1,1

# Build output array
output_arr = np.ones((m_arr , n_arr)) * np.nan

for i in range(m_arr - 2):
    for j in range(n_arr - 2):
        # Compute output pixel location for output array
        op_m, op_n = i + rpi_m, j + rpi_n

        # Get 3 x 3 magnitude slice
        arr_slice = magnitude[i:i+3, j:j+3]

        # Get 3 x 3 angle slice
        angle_slice = positive_gradient_angle[i:i+3, j:j+3]

        # If undefined value at reference pixel in magnitude or angle
        # put zero in output pixel location
        if np.isnan(arr_slice[rpi_m][rpi_n]) or
        np.isnan(angle_slice[rpi_m][rpi_n]):
            output_arr[op_m][op_n] = 0
        else:
            # Get the sector value
            sector =
            Sector().get_sector(angle_slice[rpi_m][rpi_n])

            if sector == 0:
# If undefined value at any of sector neighbour put zero in output
# pixel location
            if np.isnan(arr_slice[rpi_m][rpi_n+1]) or
np.isnan(arr_slice[rpi_m][rpi_n-1]):
                output_arr[op_m][op_n] = 0

```

```

# If reference pixel is greater than its sector neighbours put
reference pixel value at output location
    elif arr_slice[rpi_m][rpi_n] > arr_slice[rpi_m][rpi_n+1] and
arr_slice[rpi_m][rpi_n] > arr_slice[rpi_m][rpi_n-1]:
        output_arr[op_m][op_n] =
            arr_slice[rpi_m][rpi_n]

# If reference pixel value is less than its sector neighbours put
zero in output pixel location
    else:
        output_arr[op_m][op_n] = 0

    elif sector == 1:
# If undefined value at any of sector neighbour put zero in output
pixel location
        if np.isnan(arr_slice[rpi_m-1][rpi_n+1]) or
np.isnan(arr_slice[rpi_m+1][rpi_n-1]):
            output_arr[op_m][op_n] = 0

# If reference pixel is greater than its sector neighbours put
reference pixel value at output location
        elif arr_slice[rpi_m][rpi_n] > arr_slice[rpi_m-1][rpi_n+1] and
arr_slice[rpi_m][rpi_n] > arr_slice[rpi_m+1][rpi_n-1]:
            output_arr[op_m][op_n] =
arr_slice[rpi_m][rpi_n]

# If reference pixel value is less than its sector neighbours put
zero in output pixel location
        else:
            output_arr[op_m][op_n] = 0

    elif sector == 2:
# If undefined value at any of sector neighbour put zero in output
pixel location
        if np.isnan(arr_slice[rpi_m-1][rpi_n]) or
np.isnan(arr_slice[rpi_m+1][rpi_n]):
            output_arr[op_m][op_n] = 0

# If reference pixel is greater than its sector neighbours put
reference pixel value at output location
        elif arr_slice[rpi_m][rpi_n] > arr_slice[rpi_m-1][rpi_n] and
arr_slice[rpi_m][rpi_n] > arr_slice[rpi_m+1][rpi_n]:
            output_arr[op_m][op_n] =
arr_slice[rpi_m][rpi_n]

```

```

# If reference pixel value is less than its sector neighbours put
zero in output pixel location
    else:
        output_arr[op_m][op_n] = 0

    elif sector == 3:
# If undefined value at any of sector neighbour put zero in output
pixel location
    if np.isnan(arr_slice[rpi_m-1][rpi_n-1]) or
np.isnan(arr_slice[rpi_m+1][rpi_n+1]):
        output_arr[op_m][op_n] = 0

# If reference pixel is greater than its sector neighbours put
reference pixel value at output location
    elif arr_slice[rpi_m][rpi_n] > arr_slice[rpi_m-1][rpi_n-1] and
arr_slice[rpi_m][rpi_n] > arr_slice[rpi_m+1][rpi_n+1]:
        output_arr[op_m][op_n] =
arr_slice[rpi_m][rpi_n]

# If reference pixel value is less than its sector neighbours put
zero in output pixel location
    else:
        output_arr[op_m][op_n] = 0

# If sector value is other 0,1,2,3 raise an error.(Not going to
happen its there for correctness)
    else:
        raise f"Undefined sector: {sector}"
        cv2.imwrite(os.path.join(args.output_folder, image_name +
'_non_maxima_supression.bmp'), output_arr)
        return output_arr

```

5. Source Code for Thresholding:- Thresholding.py

```

import os
import numpy as np
import cv2

def perform_thresholding(args, image_name, image):
    """
    Args:
        image: Non maxima suppressed
    Returns:
        img1 : Image after applying threshold t1
    """

```

```

        img2 : Image after applying threshold t2
        img3 : Image after applying threshold t3
    ...

# Store all the values of image after non- maxima suppression
which are greater than zero into array
image_arr = image[image>0].ravel()

# Get 25th percentile of the array
t1 = np.percentile(image_arr,25)
image_1 = (image > t1).astype("int32")
cv2.imwrite(os.path.join(args.output_folder, image_name +
f'_threshold_t1_{np.round(t1, 2)}.bmp'), image_1 * 255)
# Multiplying the image with 255 for contrast

# Get 50th percentile of the array
t2 = np.percentile(image_arr,50)
image_2 = (image > t2).astype("int32")
cv2.imwrite(os.path.join(args.output_folder, image_name +
f'_threshold_t2_{np.round(t2, 2)}.bmp'), image_2 * 255)

# Get 75th percentile of the array
t3 = np.percentile(image_arr,75)
image_3 = (image > t3).astype("int32")
cv2.imwrite(os.path.join(args.output_folder, image_name +
f'_threshold_t3_{np.round(t3, 2)}.bmp'), image_3 * 255)

# Apply threshold to the image and convert it into integer array
return image_1, image_2, image_3

```

6. Common Functions are included as Utility Class :- utility.py

```

import numpy as np

# A class to store all operators
class Operator:
    # Prewitt operator for Gx
    gx = np.array([
        [-1,0,1],
        [-1,0,1],
        [-1,0,1]])

    # Prewitt operator for Gy
    gy = np.array([

```



```

        [1,1,1],
        [0,0,0],
        [-1,-1,-1]])

# Gaussian mask
gaussian_mask = np.array([
    [1,1,2,2,2,1,1],
    [1,2,2,4,2,2,1],
    [2,2,4,8,4,2,2],
    [2,4,8,16,8,4,2],
    [2,2,4,8,4,2,2],
    [1,2,2,4,2,2,1],
    [1,1,2,2,2,1,1]])

# A class to store sector angle definitions and method to provide
sector based on angle
class Sector():
    def __init__(self):
        # Dictionary with {sector: sector range}
        self.sector = {0: [(0, 22.5), (337.5, 360), (157.5, 202.5)], 1:
[(22.5, 67.5), (202.5, 247.5)], 2: [(67.5, 112.5), (247.5, 292.5)],
3: [(112.5, 157.5), (292.5, 337.5)]}

    def get_sector(self, angle):
        for key, val in self.sector.items():
            for l, u in val:
                # check if angle lies in the range if yes return key
                if angle >= l and angle < u:
                    return key

            # If angle is not in any range we return -1. (Not going to
            happen. Its there for correctness)
        return -1

# A function to apply discrete convolutions
def apply_discrete_convolution(image, mask):
    """
    Args:
        image : An image to use for convolution
        mask : An mask to use for convolution
    Returns:
        convolved image: An image after convolution
    """
    # Get the shape of image and mask
    (m_image, n_image), (m_mask, n_mask) = image.shape, mask.shape

```

```

        # Compute the reference pixel index from where output array will
        start populating
        rpi_m, rpi_n = int(np.floor(m_mask/2)), int(np.floor(n_mask/2))

        # Initialize an output array with nan values
        output_arr = np.ones((m_image, n_image)) * np.nan

        # Iterate through the image
        for i in range(m_image - m_mask + 1):
            for j in range(n_image - n_mask + 1):
                # Isolate the image slice to apply convolution
                img_slice = image[i:i+m_mask, j:j+n_mask]
                # Apply convolution and store the result in output array
                in appropriate location
                output_arr[i+rpi_m][j+rpi_n] = np.sum(img_slice * mask)

        return output_arr

# A function to convert negative angles to positive angles
def get_positive_angle(angle):
    pos_angle = angle.copy()
    pos_angle[pos_angle<0] += 360
    return pos_angle

```

Output Images

Image 1

1. Normalized image result after Gaussian smoothing

[input image – house.bmp, output image - House_gaussian_smooth_normalized.bmp]



2. **Normalized horizontal and vertical gradient responses (two separate images.)** To generate normalized gradient responses, take the absolute value of the results first and then normalize.


Gx:

[input image – house.bmp, output image – House_Gx_Normalized.bmp]



Gy:

[input image – house.bmp, output image - House_Gy_Normalized.bmp]



3. Normalized gradient magnitude image.

[input image – house.bmp, output image - House_gradient_magnitude_normalized.bmp]



4. Normalized gradient magnitude image after non-maxima suppression.


[input image – house.bmp, output image - House_non_maxima_supression.bmp]



(5) Binary edge maps using simple thresholding for thresholds chosen at the 25th, 50th and 75th percentiles

Thresholding T1:

[input image – house.bmp, output image - House_threshold_t1_2.28.bmp]



Thresholding T2:

[input image – house.bmp, output image - House_threshold_t2_5.3.bmp]



Thresholding T3:

[input image – house.bmp, output image - House_threshold_t3_15.56.bmp]



Image 2

(1) Normalized image result after Gaussian smoothing

[input image – Test patterns.bmp, output image - Test patterns_non_maxima_supression.bmp]



- (2) **Normalized horizontal and vertical gradient responses (two separate images.)** To generate normalized gradient responses, take the absolute value of the results first and then normalize.

Gx:

[input image – Test patterns.bmp, output image - Test patterns_non_maxima_supression.bmp]



Gy:

[input image – Test patterns.bmp, output image - Test patterns_non_maxima_supression.bmp]




3. Normalized gradient magnitude image.

[input image – Test patterns.bmp, output image - Test patterns_maxima_supression.bmp]



4. Normalized gradient magnitude image after non-maxima suppression.


[input image – Test patterns.bmp, output image - Test patterns_non_maxima_supression.bmp]



(5) Binary edge maps using simple thresholding for thresholds chosen at the 25th, 50th and 75th percentiles

Thresholding T1:

[input image – Test patterns.bmp, output image - Test patterns_threshold_t1_2.28.bmp]



Thresholding T2:

[input image – Test patterns.bmp, output image – Test_patterns_threshold_t2_24.45.bmp]



Thresholding T3:

[input image – Test patterns.bmp, output image - Test
patterns_threshold_t3_15.56.bmp]

